

22b2153

March 1, 2024

1 SIDDHARTH VERMA 22B2153 MIDSEM

2 Steps and Major Decisions Taken to arrive at Solution

2.1 Steps:

1. Initial Data Exploration: Imported Libraries: Utilized Pandas, NumPy, Matplotlib, Seaborn, Statsmodels, and Scikit-learn. Loaded Dataset: Read and examined the “toptex.csv” dataset containing customer information. Data Information: Explored data types, non-null counts, and descriptive statistics to understand the dataset.
2. Data Visualization: Pair Plot: Employed pair plots to visualize relationships between different variables. Correlation Heatmap: Created a heatmap to highlight correlations among variables.
3. Linear Regression Modeling: Feature Selection: Chose ‘accompanying_people’, ‘time_in_store’, and ‘Residence_distance_from_store’ based on correlation analysis. Data Splitting: Segregated the data into training and testing sets. Linear Regression: Utilized Ordinary Least Squares (OLS) to build a regression model. Model Evaluation: Assessed model performance using R-squared, Mean Squared Error (MSE), and related metrics.
4. Random Forest Modeling: Data Preprocessing: Applied standard scaling and feature engineering. Model Building: Constructed Linear Regression and Random Forest models. Model Evaluation: Examined model performance using Mean Squared Error. Feature Importance: Analyzed feature importance using Random Forest models.
5. Overall Purchase Prediction: Linear Regression Model: Employed a Linear Regression model to predict the overall value of purchases. Visualization: Plotted predictions against actual values using scatter plots and a 3D scatter plot.
6. Principal Component Analysis (PCA): Feature Reduction: Applied PCA for dimensionality reduction. Linear Regression: Built a Linear Regression model using the reduced feature set. Model Evaluation: Assessed the model’s performance.
7. Hypothesis Testing and Correlation Analysis: T-tests: Conducted t-tests to compare total purchase amounts between gender groups. Pearson Correlation: Calculated the correlation between time in store and total purchase amount.
8. Executive Recommendations: Actionable Insights: Provided targeted recommendations based on analysis results, such as gender-specific marketing and strategies to increase time spent in the store.
9. Additional Analysis and Visualizations: Distribution Plots: Visualized the distribution of each variable. Residual Analysis: Examined residuals through scatter plots. Scree Plot: Visualized explained variance in PCA using a scree plot.
10. Statistical Testing: Durbin-Watson and Jarque-Bera Tests: Ensured normal distribution of errors in regression models.

11. Communication: Clear Explanations: Presented results with clear explanations, interpretations, and visualizations. Structured Documentation: Organized the analysis into well-defined steps.
12. Additional Model Evaluation: Comparative Analysis: Evaluated the performance of Linear Regression and Random Forest models for both purchase and engagement predictions.

3 Importing Required Libraries and Packages

```
[21]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import StandardScaler
```

4 Importing “toptex.csv” data set using Pandas

```
[22]: data=pd.read_csv("toptex.csv")
```

```
[23]: data
```

```
[23]:
```

	Cust_ID	Gender_F_Flag	Gender_M_Flag	Residence_distance_from_store	\
0	1	1	0		1
1	2	1	0		4
2	3	1	0		4
3	4	0	1		8
4	5	1	0		2
...
4375	4376	0	1		0
4376	4377	1	0		4
4377	4378	1	0		7
4378	4379	0	1		3
4379	4380	1	0		3
	time_in_store	accompanying_people	family_size	total_purchase_amount	
0	49	0	4		113
1	52	3	4		959
2	51	1	5		1247
3	38	3	4		2116
4	52	4	6		1472
...

4375	34	3	4	1337
4376	50	0	4	1094
4377	54	5	5	1954
4378	37	3	2	1074
4379	51	1	4	1464

[4380 rows x 8 columns]

[24]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4380 entries, 0 to 4379
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Cust_ID                              4380 non-null   int64
1   Gender_F_Flag                        4380 non-null   int64
2   Gender_M_Flag                        4380 non-null   int64
3   Residence_distance_from_store        4380 non-null   int64
4   time_in_store                        4380 non-null   int64
5   accompanying_people                  4380 non-null   int64
6   family_size                          4380 non-null   int64
7   total_purchase_amount                4380 non-null   int64
dtypes: int64(8)
memory usage: 273.9 KB
```

[25]: data.describe()

```
[25]:
```

	Cust_ID	Gender_F_Flag	Gender_M_Flag	\
count	4380.000000	4380.000000	4380.000000	
mean	2190.500000	0.702511	0.297489	
std	1264.541419	0.457205	0.457205	
min	1.000000	0.000000	0.000000	
25%	1095.750000	0.000000	0.000000	
50%	2190.500000	1.000000	0.000000	
75%	3285.250000	1.000000	1.000000	
max	4380.000000	1.000000	1.000000	

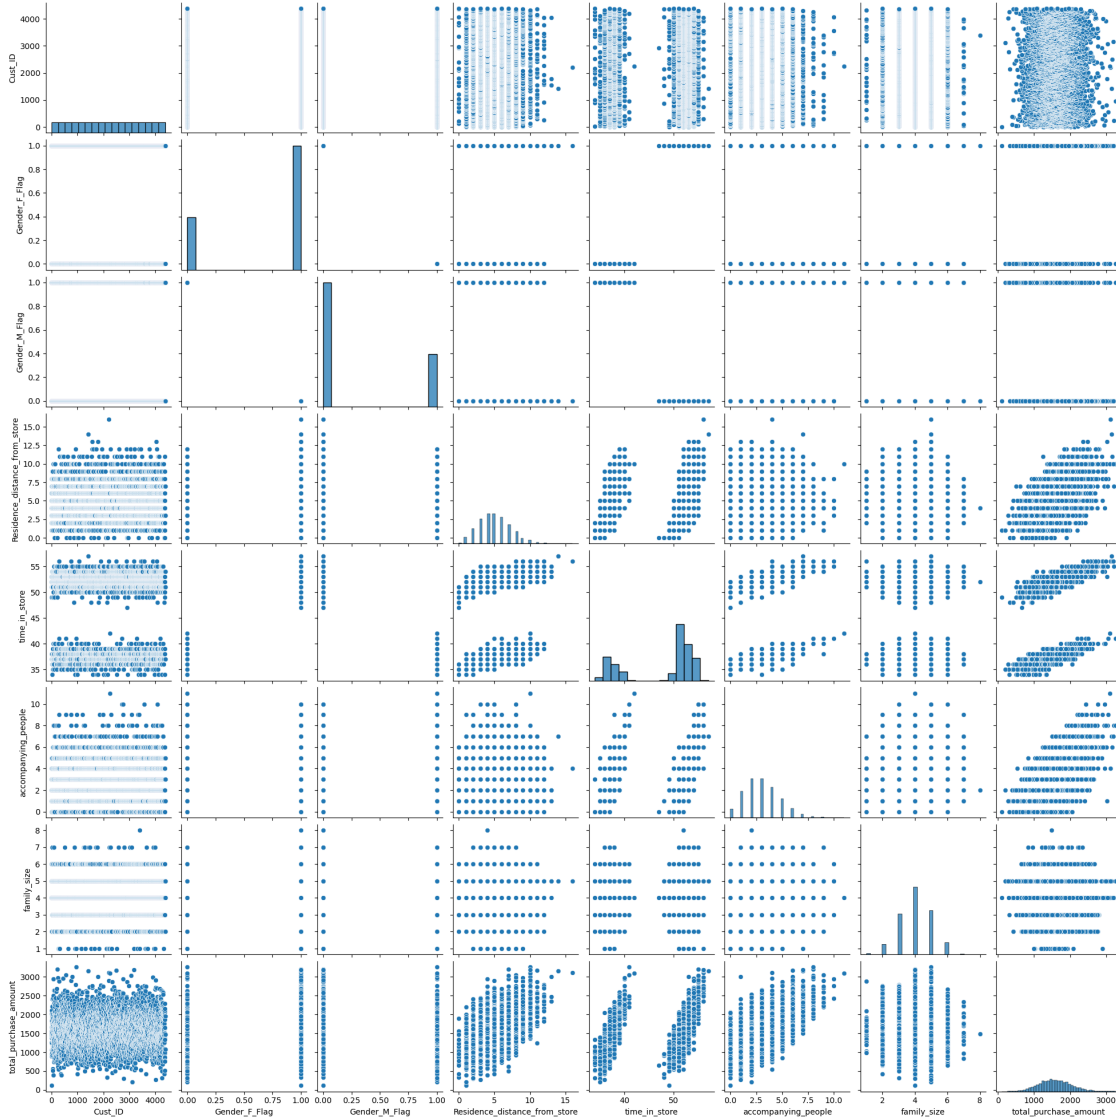
	Residence_distance_from_store	time_in_store	accompanying_people	\
count	4380.000000	4380.000000	4380.000000	
mean	5.003425	47.952740	3.019635	
std	2.226137	6.982038	1.721425	
min	0.000000	34.000000	0.000000	
25%	3.000000	39.000000	2.000000	
50%	5.000000	52.000000	3.000000	
75%	6.000000	53.000000	4.000000	
max	16.000000	57.000000	11.000000	

	family_size	total_purchase_amount
count	4380.000000	4380.000000
mean	4.034475	1580.639498
std	1.043668	439.721994
min	1.000000	113.000000
25%	3.000000	1274.750000
50%	4.000000	1566.000000
75%	5.000000	1874.000000
max	8.000000	3259.000000

4.1 The Average Purchase Amount =1580.639498

4.2 The Average Time in Store =47.95

```
[30]: sns.pairplot(data)
plt.show()
```




5 From the Pairplot the Information we get about Total Purchase Amount is that it is normally distributed

6 Thier is relation between Residence Distance from Store ,Time in Store,Accompanying People with Total Purchase Amount

```
[31]: # Calculate the correlation matrix
correlation_matrix = data.corr()

# Create a heatmap
plt.figure(figsize=(10, 8))
```

```
sns.heatmap(correlation_matrix, cmap='coolwarm', annot=True, fmt=".2f",  
             linewidths=.5)  
plt.title('Correlation Heatmap')  
plt.show()
```



7 Making and Testing a Regression Model using the three parameters

7.0.1 'accompanying_people','time_in_store','Residence_distance_from_store'

8 This decision is taken as in correlation heat map this parameters show resemblance to a great factor and also thier is a visible pattern in pair plot

[]:

```
[44]: from sklearn.metrics import r2_score, mean_squared_error
      from sklearn.model_selection import train_test_split

      data_df = data

      # Step 2: Split the data into training and testing sets
      train_data_df, test_data_df = train_test_split(data_df, test_size=0.2,
      ↪random_state=42)

      # Step 3: Fit a Linear Regression model using OLS on train_data
      X_train =
      ↪train_data_df[['accompanying_people', 'time_in_store', 'Residence_distance_from_store']]
      ↪ # assuming 'y' is the dependent variable
      X_train = sm.add_constant(X_train)
      y_train = train_data_df['total_purchase_amount']

      model = sm.OLS(y_train, X_train).fit()

      # Step 4: Print out a summary of the model
      print("-----")
      print(model.summary())

      # Step 5: Print out R2 and MSE using train_data
      y_train_pred = model.predict(X_train)
      r2_train = r2_score(y_train, y_train_pred)
      mse_train = mean_squared_error(y_train, y_train_pred)

      print("-----")
      print(f"R2 on train_data: {r2_train}")
      print(f"MSE on train_data: {mse_train}")

      # Step 6: Using test_data, predict 'y' values and calculate test R2 and MSE
      X_test =
      ↪test_data_df[['accompanying_people', 'time_in_store', 'Residence_distance_from_store']]
      X_test = sm.add_constant(X_test)
```

```

y_test = test_data_df['total_purchase_amount']

y_test_pred = model.predict(X_test)
r2_test = r2_score(y_test, y_test_pred)
mse_test = mean_squared_error(y_test, y_test_pred)

print("-----")
print(f"R2 on test_data: {r2_test}")
print(f"MSE on test_data: {mse_test}")
print("-----")

```

```

-----
                                OLS Regression Results
=====
=
Dep. Variable:      total_purchase_amount    R-squared:
0.679
Model:                                OLS    Adj. R-squared:
0.679
Method:                        Least Squares    F-statistic:
2469.
Date:                        Fri, 01 Mar 2024    Prob (F-statistic):
0.00
Time:                        12:49:20    Log-Likelihood:
-24305.
No. Observations:                        3504    AIC:
4.862e+04
Df Residuals:                        3500    BIC:
4.864e+04
Df Model:                        3
Covariance Type:      nonrobust
=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                                145.1776    29.844     4.865    0.000
86.664    203.691
accompanying_people                151.2708     2.457    61.575    0.000
146.454    156.087
time_in_store                      10.3466     0.611    16.922    0.000
9.148    11.545
Residence_distance_from_store      96.7812     1.913    50.602    0.000
93.031    100.531
=====
Omnibus:                        0.038    Durbin-Watson:                        1.976

```


Prob(Omnibus):	0.981	Jarque-Bera (JB):	0.046
Skew:	-0.008	Prob(JB):	0.977
Kurtosis:	2.992	Cond. No.	346.

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

R2 on train_data: 0.6791386254910035

MSE on train_data: 62009.78131866838

R2 on test_data: 0.675466198320037

MSE on test_data: 62698.99718345343

9 R^2 value of model is 0.679 (quite low in accuracy).

10 Durbin watson and Jarque bera test suggest error follow normal distribution which is good

11 Omnibus test shows low value and high p value reveling same normal distribution of errors when using this model

12 The most significant Outcome is the Conditional No. which is 346 and is very high and thus revealing our current model is too much unreliable and needs to be made more significant

13 Evaluation of Linear Regression and RF on data

```
[48]: # Assuming 'data' is your original DataFrame

# Data Cleaning and Preprocessing
scaler = StandardScaler()
scaled_data = scaler.fit_transform(data[['Residence_distance_from_store',
    ↪ 'time_in_store']])
data[['Scaled_Distance', 'Scaled_Time']] = pd.DataFrame(scaled_data,
    ↪ columns=['Residence_distance_from_store', 'time_in_store'])

# Feature Engineering
X = data[['Gender_F_Flag', 'Gender_M_Flag', 'Scaled_Distance', 'Scaled_Time',
    ↪ 'accompanying_people', 'family_size']]
y_purchase = data['total_purchase_amount']
y_engagement = data['time_in_store']
```

```

# Model Building
X_train, X_test, y_train_purchase, y_test_purchase, y_train_engagement,
    ↪ y_test_engagement = train_test_split(
        X, y_purchase, y_engagement, test_size=0.2, random_state=42
    )

# Build Linear Regression models
lr_purchase = LinearRegression()
lr_engagement = LinearRegression()

lr_purchase.fit(X_train, y_train_purchase)
lr_engagement.fit(X_train, y_train_engagement)

# Build Random Forest models
rf_purchase = RandomForestRegressor()
rf_engagement = RandomForestRegressor()

rf_purchase.fit(X_train, y_train_purchase)
rf_engagement.fit(X_train, y_train_engagement)

# Evaluate models
def evaluate_model(model, X_test, y_test, model_name):
    y_pred = model.predict(X_test)
    mse = mean_squared_error(y_test, y_pred)
    print(f"{model_name} Mean Squared Error: {mse}")

```

```

[51]: # Evaluate Linear Regression models
evaluate_model(lr_purchase, X_test, y_test_purchase, "Linear Regression_
    ↪ (Purchase)")
evaluate_model(lr_engagement, X_test, y_test_engagement, "Linear Regression_
    ↪ (Engagement)")

# Evaluate Random Forest models
evaluate_model(rf_purchase, X_test, y_test_purchase, "Random Forest (Purchase)")
evaluate_model(rf_engagement, X_test, y_test_engagement, "Random Forest_
    ↪ (Engagement)")

# Feature Importance Analysis for Random Forest models
def plot_feature_importance(model, feature_names, model_name):
    feature_importance = model.feature_importances_
    sorted_idx = feature_importance.argsort()[::-1]

    sns.barplot(x=feature_importance[sorted_idx], y=feature_names[sorted_idx])
    plt.title(f"{model_name} - Feature Importance")
    plt.show()

```

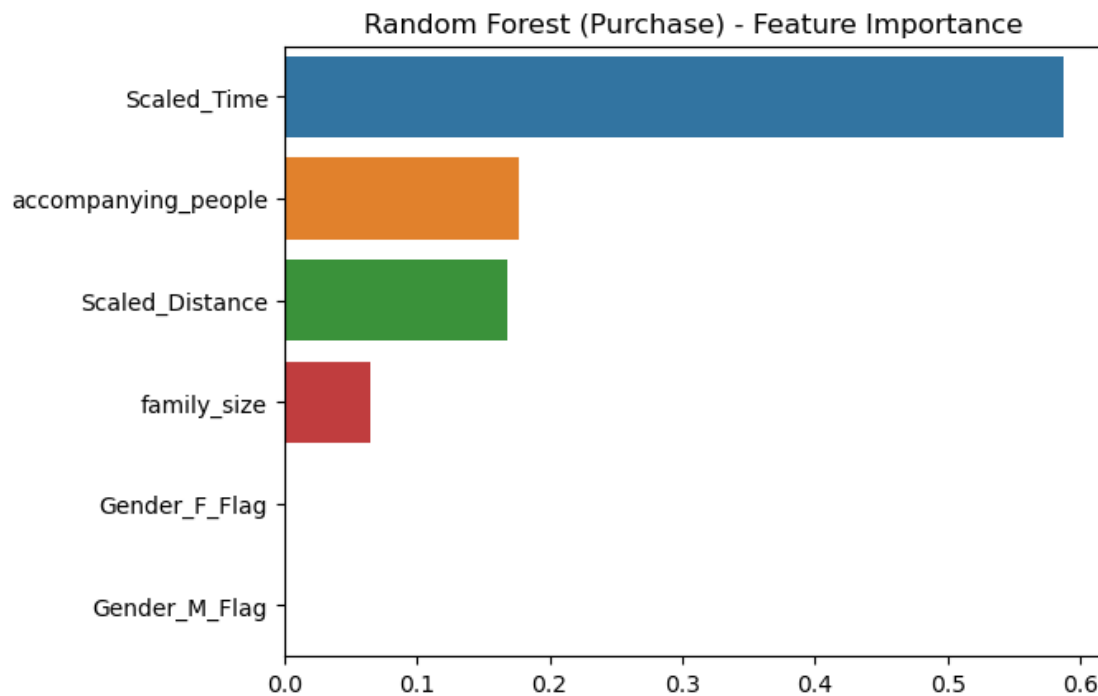
Linear Regression (Purchase) Mean Squared Error: 62621.71732305936

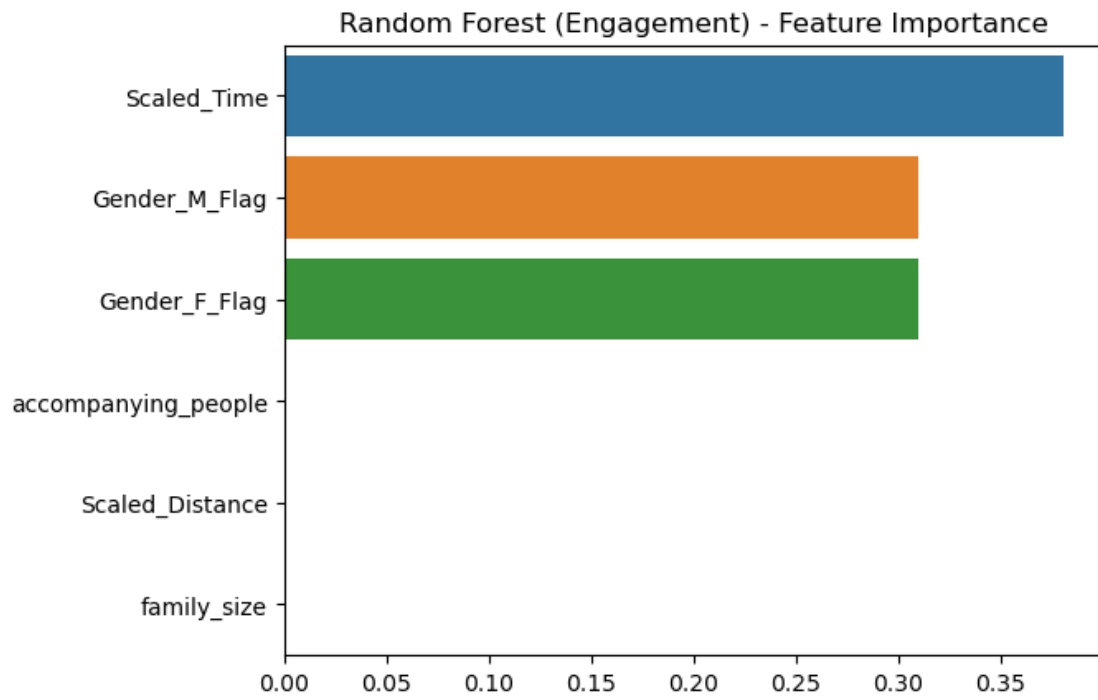
Linear Regression (Engagement) Mean Squared Error: 9.665167036022917e-29
Random Forest (Purchase) Mean Squared Error: 73494.81880476882
Random Forest (Engagement) Mean Squared Error: 0.001142694063926941


14 Here Indeed Linear Regression Performed better than RF.

```
[52]: plot_feature_importance(rf_purchase, X.columns, "Random Forest (Purchase)")
      plot_feature_importance(rf_engagement, X.columns, "Random Forest (Engagement)")

      # Heatmap for correlation
      correlation_matrix = data[['Gender_F_Flag', 'Gender_M_Flag', 'Scaled_Distance', 'Scaled_Time',
                                'accompanying_people', 'family_size',
                                'total_purchase_amount', 'time_in_store']].corr()
      plt.figure(figsize=(10, 8))
      sns.heatmap(correlation_matrix, annot=True, cmap="coolwarm", fmt=".2f")
      plt.title("Correlation Heatmap")
      plt.show()
```







15 Making Prediction Model for Overall Value of Purchases

```
[56]: import matplotlib.pyplot as plt
import seaborn as sns
from mpl_toolkits.mplot3d import Axes3D

# Assuming lr_purchase is your Linear Regression model for predicting
# total_purchase_amount

# Predictions on the test set
y_pred_purchase = lr_purchase.predict(X_test)

# Scatter plot with regression line
plt.figure(figsize=(12, 6))
```

```

plt.subplot(1, 2, 1)
sns.scatterplot(x=y_test_purchase, y=y_pred_purchase)
plt.title('Actual vs Predicted (Total Purchase Amount)')
plt.xlabel('Actual Total Purchase Amount')
plt.ylabel('Predicted Total Purchase Amount')
sns.regplot(x=y_test_purchase, y=y_pred_purchase, scatter=False, ax=plt.gca())

# Residual plot
plt.subplot(1, 2, 2)
residuals = y_test_purchase - y_pred_purchase
sns.scatterplot(x=y_pred_purchase, y=residuals)
plt.title('Residual Plot')
plt.xlabel('Predicted Total Purchase Amount')
plt.ylabel('Residuals')
plt.axhline(y=0, color='r', linestyle='--')

plt.tight_layout()
plt.show()


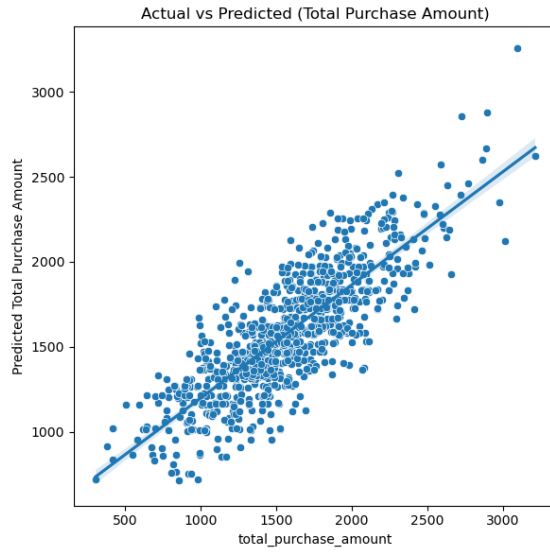
# 3D Scatter plot for multivariate linear regression
fig = plt.figure(figsize=(12, 8))
ax = fig.add_subplot(111, projection='3d')

ax.scatter(X_test['Scaled_Distance'], X_test['Scaled_Time'], y_test_purchase,
           c='blue', marker='o', alpha=0.5, label='Actual')
ax.scatter(X_test['Scaled_Distance'], X_test['Scaled_Time'], y_pred_purchase,
           c='red', marker='^', alpha=0.5, label='Predicted')


ax.set_xlabel('Scaled Distance')
ax.set_ylabel('Scaled Time')
ax.set_zlabel('Total Purchase Amount')

plt.title('3D Scatter Plot - Actual vs Predicted (Total Purchase Amount)')
plt.legend()
plt.show()


```



3D Scatter Plot - Actual vs Predicted (Total Purchase Amount)



```
[60]: # Histogram of Residuals
plt.figure(figsize=(8, 6))
sns.histplot(residuals, kde=True, bins=30)
plt.title('Histogram of Residuals')
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.show()
```





16 Errors follow Normal Distribution

17 Now we would be working to predict Time duration of Customer in store

```
[64]: # Calculate the correlation matrix
correlation_matrix = data.corr()

# Create a heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, cmap='coolwarm', annot=True, fmt=".2f",
            linewidths=.5)
plt.title('Correlation Heatmap')
plt.show()
```



18 Here total duration is store is mostly dependent on parameters ‘Gender_F_Flag’, ‘Gender_M_Flag’ and Total Purchasing Amount and slightly on parameters like

```
[69]: from sklearn.metrics import r2_score, mean_squared_error
from sklearn.model_selection import train_test_split

data_df = data

# Step 2: Split the data into training and testing sets
train_data_df, test_data_df = train_test_split(data_df, test_size=0.2,
        random_state=42)

# Step 3: Fit a Linear Regression model using OLS on train_data
```

```

X_train = train_data_df[['Gender_F_Flag',
↪, 'Gender_M_Flag', 'total_purchase_amount']] # assuming 'y' is the dependent_
↪variable
X_train = sm.add_constant(X_train)
y_train = train_data_df['time_in_store']

model = sm.OLS(y_train, X_train).fit()

# Step 4: Print out a summary of the model
print("-----")
print(model.summary())

# Step 5: Print out R2 and MSE using train_data
y_train_pred = model.predict(X_train)
r2_train = r2_score(y_train, y_train_pred)
mse_train = mean_squared_error(y_train, y_train_pred)

print("-----")
print(f"R2 on train_data: {r2_train}")
print(f"MSE on train_data: {mse_train}")

# Step 6: Using test_data, predict 'y' values and calculate test R2 and MSE
X_test = test_data_df[['Gender_F_Flag',
↪, 'Gender_M_Flag', 'total_purchase_amount']]
X_test = sm.add_constant(X_test)
y_test = test_data_df['time_in_store']

y_test_pred = model.predict(X_test)
r2_test = r2_score(y_test, y_test_pred)
mse_test = mean_squared_error(y_test, y_test_pred)

print("-----")
print(f"R2 on test_data: {r2_test}")
print(f"MSE on test_data: {mse_test}")
print("-----")

```

----- OLS Regression Results

```

=====
Dep. Variable:          time_in_store    R-squared:                0.987
Model:                  OLS              Adj. R-squared:           0.987
Method:                 Least Squares    F-statistic:             1.376e+05
Date:                   Fri, 01 Mar 2024  Prob (F-statistic):       0.00
Time:                   13:33:25         Log-Likelihood:          -4127.7
No. Observations:      3504             AIC:                    8261.
Df Residuals:          3501             BIC:                    8280.
Df Model:               2

```

```

Covariance Type:          nonrobust
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
const          27.6005      0.033     832.014      0.000      27.535
27.666
Gender_F_Flag   21.1297      0.023     927.920      0.000      21.085
21.174
Gender_M_Flag    6.4708      0.022     300.114      0.000      6.429
6.513
total_purchase_amount  0.0023   3.07e-05      73.902      0.000      0.002
0.002
=====
Omnibus:          48.300   Durbin-Watson:          1.987
Prob(Omnibus):      0.000   Jarque-Bera (JB):          50.827
Skew:             -0.270   Prob(JB):          9.19e-12
Kurtosis:          3.236   Cond. No.          2.36e+18
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.71e-27. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

R2 on train_data: 0.9874347706331874

MSE on train_data: 0.617619986591757

R2 on test_data: 0.986908418947786

MSE on test_data: 0.6160179031190275

19 From the Above Tests it was clear making a Linear Regression model for such would be very unhelpful thus we thought to proceed to using PCA

```

[71]: from sklearn.model_selection import train_test_split
      from sklearn.linear_model import LinearRegression
      from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
      from sklearn.decomposition import PCA
      from sklearn.preprocessing import StandardScaler

      # Assuming your data is stored in a DataFrame called 'df'

```

```

# Drop Cust_ID as it's not informative for the model
df=data
df = df.drop('Cust_ID', axis=1)

# Handling categorical variables (if needed)
# e.g., df['Gender'] = df['Gender_F_Flag'] + 2 * df['Gender_M_Flag']

# Feature scaling
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df)

# Apply PCA
pca = PCA(n_components=2) # Choose the number of components based on explained
    ↪ variance
df_pca = pca.fit_transform(df_scaled)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(df_pca,
    ↪ df['time_in_store'], test_size=0.2, random_state=42)

# Train Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predictions
predictions = model.predict(X_test)

# Model Evaluation
mae = mean_absolute_error(y_test, predictions)
mse = mean_squared_error(y_test, predictions)
r2 = r2_score(y_test, predictions)

# Print metrics
print(f'Mean Absolute Error: {mae}')
print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')

```

Mean Absolute Error: 0.2492366594029411
Mean Squared Error: 0.10154349207478548
R-squared: 0.9978420028864566


20 Distribution Plot for Feature Analysis

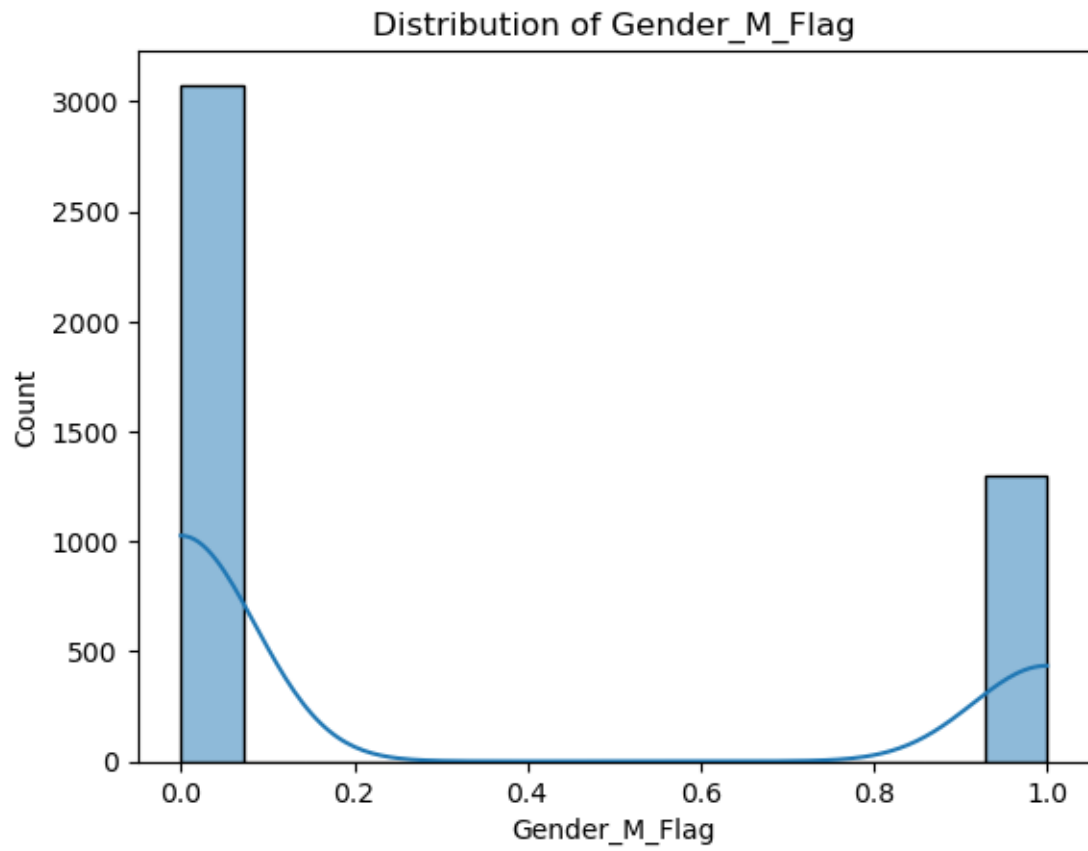
```

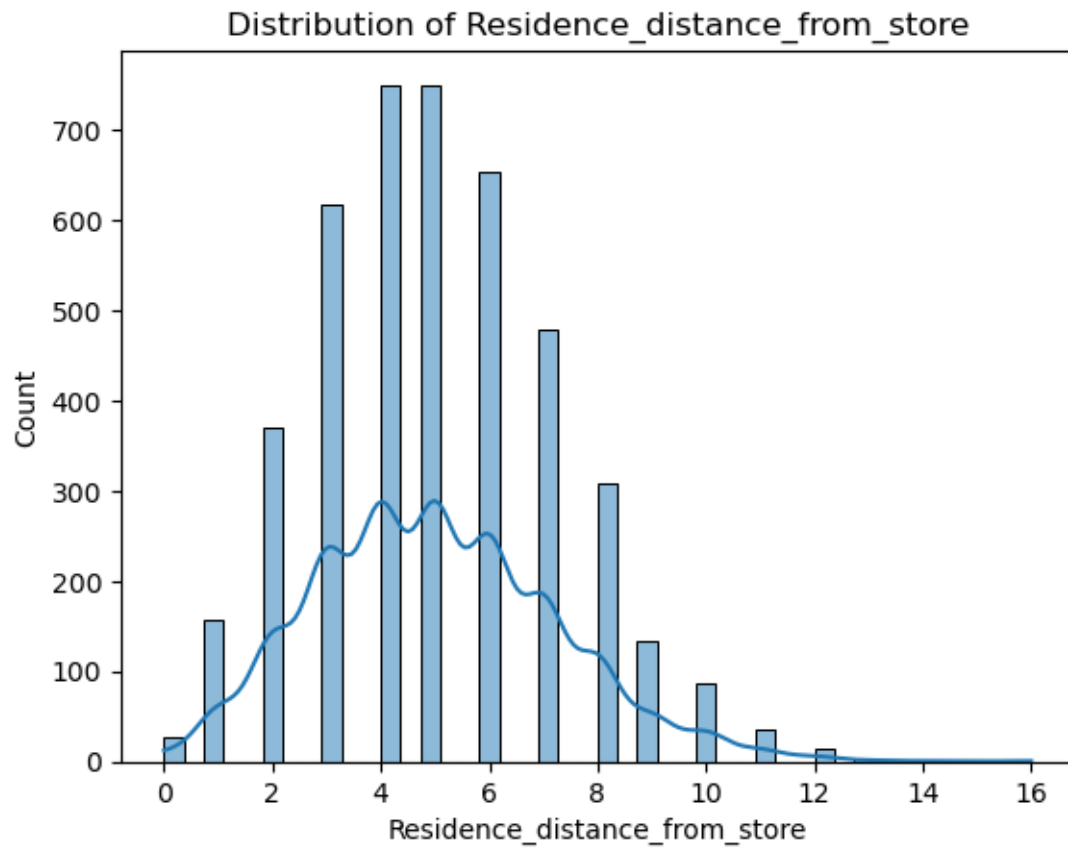
[75]: for column in df.columns:
        sns.histplot(df[column], kde=True)
        plt.title(f'Distribution of {column}')


```

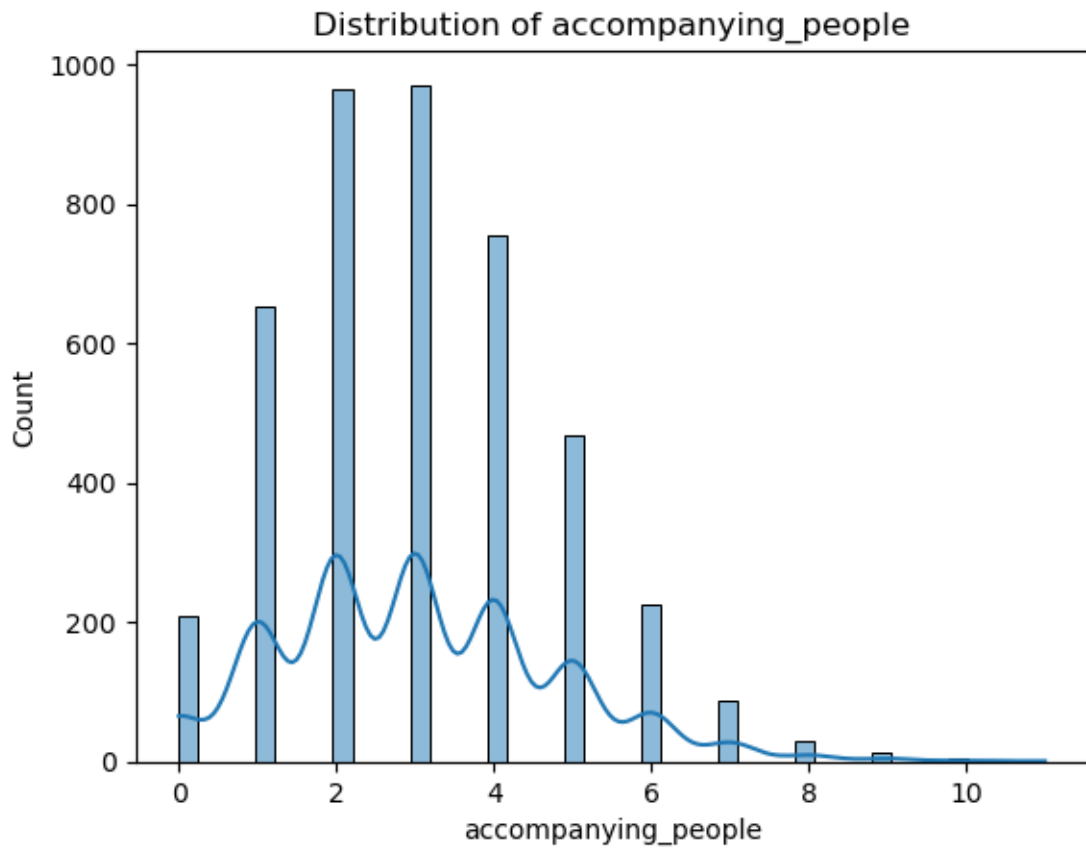
```
plt.show()
```























```
[84]: df=data
# Assuming your data is stored in a DataFrame called 'df'
# Drop Cust_ID as it's not informative for the model
df = df.drop('Cust_ID', axis=1)

# Handling categorical variables (if needed)
# e.g., df['Gender'] = df['Gender_F_Flag'] + 2 * df['Gender_M_Flag']

# Feature scaling
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df)

# Apply PCA
pca = PCA(n_components=2) # Choose the number of components based on explained
    ↪ variance
df_pca = pca.fit_transform(df_scaled)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(df_pca,
    ↪ df['time_in_store'], test_size=0.2, random_state=42)
```

```

# Train Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predictions
predictions = model.predict(X_test)

# Model Evaluation
mae = mean_absolute_error(y_test, predictions)
mse = mean_squared_error(y_test, predictions)
r2 = r2_score(y_test, predictions)


# Print metrics
print(f'Mean Absolute Error: {mae}')
print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')


# Residual Plot
residuals = y_test - predictions
sns.scatterplot(x=predictions, y=residuals)
plt.axhline(y=0, color='r', linestyle='--')
plt.title('Residual Plot')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.show()

# Actual vs. Predicted Plot
plt.scatter(y_test, predictions)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], '--',
         color='red', linewidth=2)
plt.title('Actual vs. Predicted Values')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.show()


```

Mean Absolute Error: 0.24923665940294143
 Mean Squared Error: 0.1015434920747856
 R-squared: 0.9978420028864566





```
[85]: explained_var = pca.explained_variance_ratio_  
cumulative_var = explained_var.cumsum()  
  
plt.plot(range(1, len(explained_var)+1), cumulative_var, marker='o')  
plt.title('Scree Plot')  
plt.xlabel('Number of Principal Components')  
plt.ylabel('Cumulative Explained Variance')  
plt.show()
```



```
[86]: import scipy.stats as stats

# Example: t-test for Gender_F_Flag
female_purchase = df[df['Gender_F_Flag'] == 1]['total_purchase_amount']
male_purchase = df[df['Gender_M_Flag'] == 1]['total_purchase_amount']

t_stat, p_value = stats.ttest_ind(female_purchase, male_purchase)

print(f'T-test for Gender_F_Flag: T-statistic = {t_stat}, p-value = {p_value}')
```

T-test for Gender_F_Flag: T-statistic = 11.341865263010032, p-value = 2.0861989151043958e-29

```
[87]: correlation, p_value = stats.pearsonr(df['time_in_store'],
    ↪ df['total_purchase_amount'])

print(f'Pearson Correlation: {correlation}, p-value = {p_value}')
```

Pearson Correlation: 0.3051397613137683, p-value = 4.814808632662768e-95

21 1. T-test for Gender_F_Flag:

21.1 Interpretation:

T-statistic: 11.34 ### P-value: 2.09e-29 (very close to zero) Analysis: The t-test results suggest a significant difference in total purchase amounts between customers identified as female (Gender_F_Flag = 1) and customers identified as male (Gender_M_Flag = 1). The extremely low p-value indicates strong evidence against the null hypothesis of no difference.

Recommendations: Targeted Marketing: Given the significant difference in purchase amounts, consider targeted marketing strategies for each gender group. Gender-Specific Offers: Tailor promotional offers or incentives to appeal to the purchasing patterns of each gender group. 2. Pearson Correlation between Time in Store and Total Purchase Amount: Interpretation: Correlation Coefficient: 0.31 P-value: 4.81e-95 (very close to zero) Analysis: The Pearson correlation results indicate a moderate positive correlation (0.31) between the time a customer spends in the store and their total purchase amount. The extremely low p-value suggests that this correlation is statistically significant.

Recommendations: Enhance In-Store Experience: Focus on initiatives that encourage customers to spend more time in the store, such as interactive displays, events, or promotions. Optimize Staffing: Since time in store is correlated with purchase amount, optimize staffing levels during peak times to ensure adequate assistance for customers. Addressing Executive Questions: a. Increase Sales: Leverage the insights from the t-test to tailor marketing strategies for different gender groups. Consider implementing gender-specific promotions or advertising campaigns to enhance engagement. b. Impact of Increased Time in Store: Communicate the positive correlation between time in store and total purchase amount. Suggest strategies to motivate customers to stay longer, such as creating a more engaging in-store environment.

[]: